

LinkedIn Profile Overhaul — Satyam Namdev

Status snapshot from your screenshots (Jun 28, 2026):

Item	Current state	Verdict
Headshot	Professional, clear face	Good
Banner	Generic city skyline at night	Replace
Headline	167/220 chars, no remote signal	Rewrite
Open to Work	ON, Recruiters only, SF / Hybrid+Remote	Good
Website link	Linked	Good
About section	Casual tone, opens with "Heyy"	Rewrite
Featured	Not visible / likely empty	Add
Recommendations	Not visible	Request
Skill Assessments	Not taken	Take
Custom URL	linkedin.com/in/spyrosigma	Good

1. Headline (rewrite)

Current (167/220 chars):

Lead Applied AI Engineer @ezAIx, USA | prev. ML Engineer @TunableLabs, USA | Working with LLMs (RAGs, Real-time Voice Agents, Multi Agent orchestration, MCP and more)

Problems:

- "and more" wastes chars and sounds vague
- No timezone/remote signal — recruiters filter by this
- "Working with LLMs (RAGs, Real-time Voice Agents...)" reads like a list dump, not a value prop

Suggested headline (copy-paste, 215/220 chars):

Lead AI Engineer @ezAIx, USA | Ex-Founding ML Engineer @TunableLabs | RAG Pipelines · Real-time Voice AI · LangGraph Agents | IIT Madras | IST, 4hr US overlap | Open to Remote

Alternative (more concise, 198/220 chars):

Lead AI Engineer @ezAIx (USA) | Built production RAG, Voice AI & LLM Agents at US startups | Django · Azure · LangGraph | IIT Madras | Open to Remote

Pick one. The key additions are: **"Open to Remote"** and **"IST, 4hr US overlap"** — these are the exact strings recruiters search for.

2. Banner Image (replace)

Your current banner is a generic night cityscape. It's wasted real estate.

What to make (use Canva — free LinkedIn banner template, 1584 x 396 px):

Design a clean, dark-themed banner with:

Layout:

Left side:	Right side:
SATYAM NAMDEV	[tech icons or subtle code graphic]
AI & Backend Engineer	
Built production AI systems at US startups, from India.	Django · Azure · LangGraph · RAG Python · FastAPI · Redis · Celery
	spyrosigma.in

Design rules:

- Dark background (charcoal/navy — matches your profile photo mood)
- Clean sans-serif font (Inter, DM Sans, or similar)
- No gradients, no neon, no stock photos
- Left-aligned text, right-aligned tech stack pills or subtle visual
- Your name LARGE on the left, role underneath, one-liner underneath
- Tech stack keywords on the right (recruiters scan these)
- Portfolio URL bottom-right corner

Canva search terms: "LinkedIn banner minimalist tech" or "LinkedIn cover developer dark"

3. About Section (rewrite)

Current:

Heyy 🙌 I am a Data Science enthusiast, currently working in Gen-AI field. I post insights on X (@Spyrosigma) How skilled I'm? - Heavily worked with Flask, FastAPI and I know how to debug a hellish-messy code...

Problems:

- "Heyy" — too casual, loses senior recruiters in the first word
- Self-deprecating ("is this a thing to worry about?") — never plant doubt
- List of skills with no context — reads like a student, not an engineer with production experience

- No quantified impact
- Only 3 lines show before "see more" — you're wasting them

Rewritten About (copy-paste ready):

I build production AI systems for US companies – real-time voice agents, RAG pipelines, and LLM orchestration – from India (IST, 4hr US–Eastern overlap, 1.5+ years fully remote).

Currently: Lead AI Engineer at ezAIx, where I architected VoiceCoreIQ – an enterprise contact center platform handling real-time AI voice calls via Azure OpenAI GPT-4o, intelligent routing across 18 Django apps, and a RAG pipeline serving production knowledge bases. Cut TTS costs by 85% and call setup latency by 70%.

Previously: Founding AI Engineer at TunableLabs (SF), where I built Tralyx – a Legal-AI platform processing 5,000+ legal documents with LangGraph-based entity extraction at 92% accuracy, serving 200+ legal professionals with 99.9% uptime.

What I work with daily:

- AI/LLM: Azure OpenAI, LangGraph, LangChain, RAG, Pinecone, vector search
- Backend: Django (ASGI), FastAPI, Celery, Redis, WebSockets, SSE
- Cloud: Azure (ACS, Event Grid, Blob, Document Intelligence), Docker
- Data: MongoDB, PostgreSQL, Supabase

BS in Data Science – IIT Madras.

Open to remote backend/AI engineering roles. Let's talk:
namdev2003satyam@gmail.com

Why this works:

- First 3 lines (visible before "see more") = who you are + what you build + remote proof
- Quantified impact in both roles
- Clean tech stack section that's scannable
- Ends with a clear CTA
- No casual filler, no self-doubt, no emojis

4. Featured Section (add these)

Go to your profile → "Add section" → "Featured"

Add in this order (first item shows largest):

1. **Portfolio website** — link to spyrosigma.in

- Title: "Portfolio — AI & Backend Projects"
 - Description: "Production systems I've built: VoiceCoreIQ, Legal-AI, MedMitra, and more."
2. **VoiceCoreIQ post** — link to your LinkedIn post about getting PPO from ezAlx
 - Already exists: https://www.linkedin.com/posts/spyrosigma_just-got-ppo-from-ezaix-joined-as-backend-activity-7381734793866219520-CBVU
 3. **Tralyx launch post** — link to your Tralyx launch post
 - Already exists: https://www.linkedin.com/posts/spyrosigma_aiforlaw-legalai-legaltechstartup-activity-7326485798726356992-JIyP
 4. **GitHub profile** — link to <https://github.com/Spyrosigma>
 - Title: "GitHub — Open Source & Side Projects"
 5. **(Future)** Once you record a Loom intro, add it here as the #1 item.
 6. **(Future)** Once you write your first blog post, add it here.

5. Experience Section — Rewrite bullets for remote signal

Your current experience descriptions are likely focused on *what* you built. Rewrite to also show *how* you worked remotely.

ezAlx Inc. — Lead AI & Software Engineer

Architected VoiceCoreIQ — a single-tenant Azure Communication Services-backed enterprise contact center platform, owning the full backend across 18 Django ASGI applications.

Key contributions:

- Built production AI Voicebot bridging Azure OpenAI GPT-4o Realtime with ACS media streams — sub-second latency with tool-calling for order lookup and CSAT scoring.
- Engineered end-to-end RAG pipeline: Azure Blob → Document Intelligence & Firecrawl → llama-text-embed-v2 → Pinecone namespaces, serving real-time knowledge bases.
- Designed visual node-based workflow engine for dynamic call routing, IVR trees, and omnichannel orchestration.
- Built real-time Agent Presence Dashboard via ASGI SSE middleware + Redis Pub/Sub for live supervisor monitoring.
- Reduced TTS costs by 85% and call setup latency by 70% via 3-tier caching strategy (Redis → Blob Storage → Generation).
- Drove async engineering across US and India timezones — daily written

standups, Loom walkthroughs, and PR-based code reviews as primary collaboration mode.

- Flew to Dubai for company kickoff (Dec 2025) – selected engineers gathered for strategic planning.

Stack: Django ASGI · Azure (OpenAI, ACS, Event Grid, Blob) · MongoDB · Redis · Celery · Pinecone · Docker

TunableLabs — Founding AI Engineer

First engineering hire. Built Tralyx – an AI-powered Legal Intelligence Platform from zero to production, end-to-end.

- Designed and shipped the full stack: FastAPI backend, Next.js frontend, Supabase (PostgreSQL) for multi-tenant chat sessions and document collections.
- Built LangGraph-based Entity Extraction Agent with schema validation and multi-agent orchestration across 5,000+ legal documents – 92% extraction accuracy.
- Implemented multi-LLM orchestration with streaming responses and real-time WebSocket updates; integrated web search via Tavily and Exa APIs.
- Refactored a tightly coupled Gradio monolith into modular FastAPI services, enabling API integration across multiple law firms.
- Served 200+ legal professionals with 99.9% uptime and 100+ concurrent users.
- Fully remote collaboration with SF-based founder – async-first workflow using GitHub, Notion, and Slack across IST/PST timezones.

Stack: FastAPI · LangGraph · Supabase · Weaviate · Next.js · Docker

Note the additions:

- "Drove async engineering across timezones" — remote signal
- "Async-first workflow using GitHub, Notion, and Slack" — tool fluency
- "Loom walkthroughs, PR-based code reviews" — async communication proof

6. Open to Work Settings (fine-tune)

From your screenshot, you have:

- **San Francisco, CA | Hybrid · Remote** — Good

But also add:

- **Job titles** (add multiple): "AI Engineer", "Backend Engineer", "ML Engineer", "Software Engineer", "Senior AI Engineer", "Senior Backend Engineer"
 - **Location types**: Make sure "Remote" is checked as primary
 - **Start date**: "Immediately" or whenever accurate
-

7. Skills to Add / Reorder

Go to Skills section. Add these if missing and **reorder so the top 3 visible ones are**:

1. **Python** (most searched, most relevant)
2. **Machine Learning** (broad, high-search-volume)
3. **Django** or **LLMs** (depending on target role)

Then add / ensure these are listed:

- FastAPI
- Azure
- LangChain
- LangGraph
- RAG (Retrieval-Augmented Generation)
- Docker
- PostgreSQL
- MongoDB
- Redis
- REST APIs
- WebSockets
- Celery
- Git

Take LinkedIn Skill Assessments for:

- Python
- Django
- Machine Learning
- Git
- REST APIs

These are free, take 15 min each, and add a "Verified" badge that boosts search ranking.

8. Recommendations to Request

Send personalized messages (not LinkedIn's default text) to these people:

Template for manager/lead at ezAlx:

```
Hi [Name],
```

Hope you're doing well! I'm updating my LinkedIn profile and was wondering if you'd be open to writing a short recommendation about our work together on VoiceCoreIQ.

Specifically, it would be great if you could mention:

- The scope of what I owned on the backend
- How I handled working async across timezones
- Any specific impact you noticed (cost reductions, system reliability, etc.)

Totally understand if you're too busy – no pressure at all. Happy to write one for you too!

Thanks,
Satyam

Template for TunableLabs founder:

Hi [Name],

I hope Tralyx is going well! I'm polishing up my LinkedIn and was wondering if you'd be willing to write a brief recommendation.

If you could touch on any of these, it would mean a lot:

- Building the platform from zero as the first engineer
- The LangGraph entity extraction pipeline and its accuracy
- Working fully remote across IST/PST

Happy to return the favor. Thanks!

Satyam

Also ask:

- A peer/colleague at ezAlx (Sambhav Seth, if comfortable)
- A professor or mentor from IIT Madras
- Anyone from a hackathon or open-source collaboration

Target: 5 recommendations minimum.

9. Activity / Posting Strategy

Your posting is sporadic. Here's a minimal sustainable plan:

Post 1x per week (pick a day, e.g., Tuesday):

Post types that work for your niche (rotate these):

1. **"Here's what I learned building X"** — Technical insight from VoiceCoreIQ or Tralix

- Example: "How I reduced TTS costs by 85% with a 3-tier caching strategy"
- Example: "What I learned building a LangGraph entity extraction agent over 5000 legal docs"

2. **"Hot take / opinion on AI tooling"** — Something you actually believe

- Example: "LangGraph > vanilla LangChain for production agent systems. Here's why."
- Example: "Most RAG tutorials skip the hardest part: chunking strategy matters more than your embedding model."

3. **"How I work remotely"** — Process posts

- Example: "I work with a US team from India. Here's my async communication setup."

4. **"Something I shipped this week"** — Small wins, screenshots, demos

Comment 3-5x per week:

- Follow: remote-first companies you'd want to work at
- Follow: engineering leaders at target companies
- Leave thoughtful 2-3 sentence comments on their posts (not "Great post!")
- This makes your name familiar before you ever apply

10. Location Field

Current: India

Change to: Either keep "India" (honest) OR if LinkedIn allows, add to your headline/about that you're based in Delhi/NCR, IST timezone.

The key remote signal isn't your country — it's your **timezone and overlap hours**. You already have this in the "Open to Work" settings pointing to SF, which is smart. The About section rewrite above covers the rest.

11. Custom URL

Current: linkedin.com/in/spyrosigma — Already good. No change needed.

12. Projects Section (add if not already there)

Go to "Add section" → "Additional" → "Projects"

Project 1: VoiceCoreIQ — Enterprise Contact Center Platform

```
Associated with: ezAIx Inc.  
Date: May 2025 – Present  
URL: (leave blank if NDA / internal)
```

Architected the full backend for an enterprise contact center platform on Azure.

18 Django ASGI applications handling real-time AI voice calls, intelligent call routing, queue orchestration, RAG-powered knowledge bases, and live agent monitoring.

Reduced TTS costs by 85% and call setup latency by 70%.

Handled production traffic with real-time Azure OpenAI GPT-4o voice integration.

Project 2: Tralyx — AI-Powered Legal Intelligence Platform

Associated with: TunableLabs, LLC

Date: Nov 2024 – Apr 2025

URL: <https://www.tralyx.com>

Built the core AI backend for a legal intelligence platform. LangGraph-based entity extraction across 5,000+ legal documents (92% accuracy), multi-LLM orchestration, streaming WebSocket responses, and web search integration.

Served 200+ legal professionals with 99.9% uptime.

Project 3: MedMitra — AI Medical Case Management

Date: [whenever you built it]

URL: [GitHub link if public]

Multi-agent system using LangGraph for medical case analysis – processing 500+ patient notes, lab reports (LlamaParse), and radiology images (Llama 4 vision) with 94% accuracy. SOAP notes and diagnostic suggestions in under 3 seconds.

Quick Checklist

- ☐ Rewrite headline (copy from Section 1)
- ☐ Replace banner (design on Canva, see Section 2)
- ☐ Rewrite About section (copy from Section 3)
- ☐ Add Featured items (Section 4)
- ☐ Rewrite Experience bullets (Section 5)
- ☐ Fine-tune Open to Work job titles (Section 6)
- ☐ Reorder skills + take assessments (Section 7)

- ☐ Send 5 recommendation requests (Section 8)
- ☐ Schedule first LinkedIn post (Section 9)
- ☐ Add Projects section (Section 12)
- ☐ Record 60-sec Loom intro (add to Featured when done)
- ☐ Write first technical blog post (add to Featured when done)